

Inspect Kubernetes Resources

Objectives

- Query, format, and filter attributes of Kubernetes resources.

Kubernetes and OpenShift Resources

Kubernetes uses API resource objects to represent the intended state of everything in the cluster. All administrative tasks require creating, viewing, and changing the API resources. Use the `oc api-resources` command to view the Kubernetes resources.

NOTE

The `oc` commands in the examples are identical to the equivalent `kubectl` commands.

```
[user@host ~]$ oc api-resources
NAME           SHORTNAMES   APIVERSION   NAMESPACED   KIND
bindings       v1           true          Binding
componentstatuses   cs          v1           false        ComponentStatus
configmaps      cm          v1           true         ConfigMap
endpoints       ep          v1           true         Endpoints
...output omitted...
daemonsets     ds          apps/v1      true         DaemonSet
deployments     deploy      apps/v1      true         Deployment
replicasets    rs          apps/v1      true         ReplicaSet
statefulsets   sts         apps/v1      true         StatefulSet
...output omitted...
```

The `SHORTNAME` for a component helps to minimize typing long CLI commands. For example, you can use `oc get cm` instead of `oc get configmaps`.

The `APIVERSION` column divides the objects into API groups. The column uses the `<API-Group>/<API-Version>` format. The API-Group object is blank for Kubernetes core resource objects.

Many Kubernetes resources exist within the context of a Kubernetes namespace. Kubernetes namespaces and OpenShift projects are broadly similar. A one-to-one relationship always exists between a namespace and an OpenShift project.

The `KIND` is the formal Kubernetes resource schema type.

The `oc api-resources` command can further filter the output with options that operate on the data.

Table 2.1. The `api-resources` Command Options

Option Example	Description
<code>--namespaced=true</code>	If false, return non-namespaced resources, otherwise return namespaced resources
<code>--api-group apps</code>	Limit to resources in the specified API group. Use <code>--api-group=''</code> to show core resources.
<code>--sort-by name</code>	If non-empty, sort the list of resources by using the specified field. The field can be either 'name' or 'kind'.

For example, use the following `oc api-resources` command to see all the namespaced resources in the `apps` API group, sorted by name.

```
[user@host ~]$ oc api-resources --namespaced=true --api-group apps --sort-by name
NAME           SHORTNAMES   APIVERSION   NAMESPACED   KIND
controllerrevisions   apps/v1      true          ControllerRevision
daemonsets     ds          apps/v1      true         DaemonSet
deployments     deploy      apps/v1      true         Deployment
replicasets    rs          apps/v1      true         ReplicaSet
statefulsets   sts         apps/v1      true         StatefulSet
```

Each resource contains fields that identify the resource or that describe the intended configuration of the resource. Use the `oc explain` command to get information about valid fields for an object. For example, execute the `oc explain pod` command to get information about possible pod object fields.

```
[user@host ~]$ oc explain pod
KIND:     Pod
VERSION:  v1

DESCRIPTION:
Pod is a collection of containers that can run on a host. This resource is
created by clients and scheduled onto hosts.

FIELDS:
apiVersion  <string>
  APIVersion defines the versioned schema of this representation of an
  ...
  ...
kind <string>
  Kind is a string value representing the REST resource this object
  ...
  ...
metadata    <ObjectMeta>
  Standard object's metadata. More info:
  ...
  ...
spec <PodSpec>
  Specification of the desired behavior of the pod. More info:
  ...output omitted...
```

Every Kubernetes resource contains the kind, apiVersion, spec, and status fields. However, when you create an object definition, you do not need to provide the status field. Instead, Kubernetes generates the status field, and it lists information such as runtime status and readiness. The status field is useful for troubleshooting an error or for verifying the current state of a resource.

You can use the YAML path to a field and dot-notation to get information about a particular field. For example, the following oc explain command shows details for the pod.spec fields.

```
[user@host ~]$ oc explain pod.spec
KIND:     Pod
VERSION:  v1

FIELD: spec <PodSpec>

DESCRIPTION:
  Specification of the desired behavior of the pod. More info:
  https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status
  PodSpec is a description of a pod.

FIELDS:
  activeDeadlineSeconds      <integer>
    Optional duration in seconds the pod may be active on the node relative to
    ...output omitted...
```

The following Kubernetes main resource types can be created and configured by using a YAML or a JSON manifest file, or by using OpenShift management tools:

Pods (pod)

Represent a collection of containers that share resources, such as IP addresses and persistent storage volumes. It is the primary unit of work for Kubernetes.

Services (svc)

Define a single IP/port combination that provides access to a pool of pods. By default, services connect clients to pods in a round-robin fashion.

ReplicaSet (rs)

Ensure that a specified number of pod replicas are running at any given time.

Persistent Volumes (pv)

Define storage areas for Kubernetes pods to use.

Persistent Volume Claims (pvc)

Represent a request for storage by a pod. PVCs link a PV to a pod so that its containers can use the provisioned storage, usually by mounting the storage into the container's file system.

ConfigMaps (cm) and Secrets

Contain a set of keys and values that other resources can use. Configuration maps and secrets centralize configuration values for several resources. Secrets differ from configuration maps in that the values of Secrets are always encoded (not encrypted), and their access is restricted to authorized users.

Deployment (deploy)

A deployment represents a set of containers that are included in a pod, and the deployment strategies to use. A deployment object contains the configuration to apply to all containers of each pod replica, such as the base image, tags, storage definitions, and the commands to execute when the containers start. Although Kubernetes replicas can be created stand-alone in OpenShift, they are usually created by higher-level resources such as deployment controllers.

Red Hat OpenShift Container Platform (RHOCP) adds the following main resource types to Kubernetes:

Build (build)

A *build* is the generic process of taking an input source, such as the source code of an application, and producing a usable resource as the output, such as a runnable image.

BuildConfig (bc)

A *BuildConfig* object defines a build process. A *BuildConfig* object requires at least the source input of the build, the strategy that defines how to build the input, and where to store the output of the process.

Routes

Represent a DNS hostname that the OpenShift router recognizes as an ingress point for applications and microservices.

Structure of Resources

Almost every Kubernetes object includes two nested object fields that govern the object's configuration: the object spec and the object status. The spec object describes the intended state of the resource, and the status object describes the current state. You specify the spec section of the resource when you create the object. Kubernetes controllers continuously update the status of the object throughout the existence of the object. The Kubernetes control plane continuously and actively manages every object's actual state to match the desired state you supplied.

The status field uses a collection of condition resource objects with the following fields.

Table 2.2. Condition Resource Fields

Field	Example	Description
Type	ContainersReady	The type of the condition
Status	False	The state of the condition
Reason	RequirementsNotMet	An optional field to provide extra information
Message	2/3 containers are running	An optional textual description for the condition
LastTransitionTime	2023-03-07T18:05:28Z	The last time that conditions were changed

For example, in Kubernetes, a Deployment object can represent an application that is running on your cluster. When you create a Deployment object, you might configure the deployment spec object to specify that you want three replicas of the application to be running. Kubernetes reads the deployment spec object and starts three instances of your chosen application, and updates the status field to match your spec object. If any of those instances fails, then Kubernetes responds to the difference between the spec and status objects by making a correction, in this case to start a replacement instance.

Other common fields provide base information in addition to the spec and status fields of a Kubernetes object.

Table 2.3. API Resource Fields

Field	Description
apiVersion	Identifier of the object schema version.
kind	Schema identifier.
metadata.name	Creates a label with a name key that other resources in Kubernetes can use to find it.
metadata.namespace	The namespace, or the RHOCP project where the resource is.
metadata.labels	Key-value pairs that can connect identifying metadata with Kubernetes objects.

Resources in Kubernetes consist of multiple objects. These objects define the intended state of a resource. When you create or modify an object, you make a persistent record of the intended state. Kubernetes reads the object and modifies the current state accordingly.

All RHOCP and Kubernetes objects can be represented as a JSON or YAML structure. Consider the following pod object in the YAML format:

```

apiVersion: v1 1
kind: Pod 2
metadata: 3
  name: wildfly 4
  namespace: my_app 5
  labels:
    name: wildfly 6
spec: 7
  containers:
    - resources:
        limits:
          cpu: 0.5
      image: quay.io/example/todojee:v1 8
      name: wildfly 9
      ports:
        - containerPort: 8080 10
          name: wildfly
    env: 11
      - name: MYSQL_DATABASE
        value: items
      - name: MYSQL_USER
        value: user1
      - name: MYSQL_PASSWORD
        value: mypass
...object omitted...
status: 12
  conditions:
    - lastProbeTime: null
      lastTransitionTime: "2023-08-19T12:59:22Z"
      status: "True"
      type: PodScheduled

```

- 1** Identifier of the object schema version.
- 2** Schema identifier. In this example, the object conforms to the pod schema.
- 3** Metadata for a given resource, such as annotations, labels, name, and namespace.
- 4** A unique name for a pod in Kubernetes that enables administrators to run commands on it.
- 5** The namespace, or the RHOCP project that the resource resides in.
- 6** Creates a label with a name key that other resources in Kubernetes, usually a service, can use to find it.
- 7** Defines the pod object configuration, or the intended state of the resource.
- 8** Defines the container image name.
- 9** Name of the container inside a pod. Container names are important for oc commands when a pod contains multiple containers.
- 10** A container-dependent attribute to identify the port that the container uses.
- 11** Defines a collection of environment variables.
- 12** Current state of the object. Kubernetes provides this field, which lists information such as runtime status, readiness, and container images.

Labels are key-value pairs that you define in the `.metadata.labels` object path, for example:

```

kind: Pod
apiVersion: v1
metadata:
  name: example-pod
  labels:
    app: example-pod
    group: developers
...object omitted...

```

The preceding example contains the app=example-pod and group=developers labels. Developers often use labels to target a set of objects by using the -l or the --selector option. For example, the following oc get command lists pods that contain the group=developers label:

```
[user@host ~]$ oc get pod --selector group=developers
NAME           READY   STATUS    RESTARTS   AGE
example-pod-6c9f758574-7fhg   1/1     Running   5          11d
```

Command Outputs

The kubectl and oc CLI commands provide many output formatting options. By default, many commands display a small subset of the most useful fields for the given resource type in a tabular output. Many commands support a -o wide option that shows additional fields.

Table 2.4. Tabular Fields

oc get pods	oc get pods -o wide	Example value
NAME	NAME	example-pod
READY	READY	1/1
STATUS	STATUS	Running
RESTARTS	RESTARTS	5
AGE	AGE	11d
	IP	10.8.0.60
	NODE	master01
	NOMINATED NODE	<none>
	READINESS GATES	<none>

To view all the fields that are associated with a resource, the describe subcommand shows a detailed description of the selected resource and related resources. You can select a single object by name, or all objects of that type, or provide a name prefix, or a label selector.

For example, the following command first looks for an exact match on the TYPE object and the NAME-PREFIX object. If no such resource exists, then the command outputs details for every resource of that type with a name with a NAME_PREFIX prefix.

```
[user@host ~]$ oc describe TYPE NAME-PREFIX
```

The describe subcommand provides detailed human-readable output. However, the format of the describe output might change between versions, and thus is not recommended for script development. Any scripts that rely on the output of the describe subcommand might break after a version update.

Kubernetes provides YAML and JSON-formatted output options that are suitable for parsing or scripting.

YAML Output

The -o yaml option provides a YAML-formatted output that is parsable and still human-readable.

```
[user@host ~]$ oc get pods -o yaml
apiVersion: v1
items:
- apiVersion: v1
  kind: Pod
  metadata:
    annotations:
...object omitted...
```

NOTE

The reference documentation provides a more detailed introduction to YAML.

You can use any tool that can process YAML documents to filter the YAML output for your chosen field. For example, you can use the `yq` tool at <https://mikefarah.gitbook.io/yq/> to process YAML and JSON files.

The `yq` processor uses a dot notation to separate field names in a query. The following example pipes the YAML output to the `yq` command to parse the `podIP` field.

```
[user@host ~]$ oc get pods -o yaml | yq '.items[0].status.podIP'
10.8.0.60
```

The `[0]` in the example specifies the first index in the `items` array.

NOTE

Another tool named `yq` is at <https://kislyuk.github.io/yq/>. The two `yq` tools are not compatible; commands that are designed for one of them do not work with the other.

JSON Output

Kubernetes uses the JSON format internally to process resource objects. Use the `-o json` option to view a resource in the JSON format.

```
[user@host ~]$ oc get pods -o json
{
  "apiVersion": "v1",
  "items": [
    {
      "apiVersion": "v1",
      "kind": "Pod",
      "metadata": {
        "annotations": {
```

You can use other tools to process JSON documents, such as the `jq` tool at <https://jqlang.github.io/jq/>. Similar to the `yq` processor, use the `jq` processor and dot notation on the fields to query specific information from the JSON-formatted output.

```
[user@host ~]$ oc get pods -o json | jq '.items[0].status.podIP'
"10.8.0.60"
```

Alternatively, the example might have used `.items[].status.podIP` for the query string. The empty brackets instruct the `jq` tool to query all items.

Custom Output

Kubernetes provides a custom output format that combines the convenience of extracting data via `jq` styled queries with a tabular output format. Use the `-o custom-columns` option with comma-separated `<column name>:<jq query string>` pairs.

```
[user@host ~]$ oc get pods \
-o custom-columns=PodName:".metadata.name", \
ContainerName:"spec.containers[].name", \
Phase:"status.phase", \
IP:"status.podIP", \
Ports:"spec.containers[].ports[].containerPort"
PodName          ContainerName  Phase     IP           Ports
myapp-77fb5cd997-xplhz  myapp       Running  10.8.0.60  <none>
```

Kubernetes also supports the use of JSONPath expressions. JSONPath is a query language for JSON. JSONPath expressions refer to a JSON data structure; they filter and extract formatted fields from JSON objects.

In the following example, the JSONPath expression uses the range operator to iterate over the list of pods to extract the name of the pod, its IP address, and the assigned ports.

```
[user@host ~]$ oc get pods \
-o jsonpath='{range .items[]}{\"Pod Name: \"}{.metadata.name}
{\"IP: \"}{.status.podIP}
{\"Ports: \"}{.spec.containers[].ports[].containerPort}{\"\\n\"}{end}'
Pod Name: myapp-77fb5cd997-xplhz
IP: 10.8.0.60
Ports:
```

You can customize the format of the output with Go templates, which the Go programming language uses. Use the `-o go-template` option followed by a Go template, where Go expressions are inside double braces, `{{ }}`.

```
[user@host ~]$ oc get pods \
-o go-template='{{range .items}}{{.metadata.name}}\n{{end}}'
myapp-77fb5cd997-xplhz
```

REFERENCES

For more information, refer to the *OpenShift CLI (oc)* chapter in the Red Hat OpenShift Container Platform 4.18 *CLI Tools* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html-single/cli_tools/index

For more information about custom columns, refer to the *oc get* section in the Red Hat OpenShift Container Platform 4.18 *CLI Tools* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html-single/cli_tools/index#oc-get

For more detailed introduction to YAML, see the *YAML in a Nutshell* chapter in the Red Hat Enterprise Linux Atomic Host 7 *Getting Started with Kubernetes* documentation at https://docs.redhat.com/en/documentation/red_hat_enterprise_linux_atomic_host/7/html-single/getting_started_with_kubernetes/index#overview_3

Labels and selector details are available in the *Working with Kubernetes Objects* section of the [Kubernetes Documentation - Labels and Selectors](#)

Kubernetes object details are available in the *Understanding Kubernetes Objects* section of the [Kubernetes Documentation - Understanding Kubernetes Objects](#)

Command output details are available in the *Get* section of the [Kubernetes Documentation - Getting Started](#)

For more information on JSONPath operators and syntax, see the [Kubernetes Documentation - JSONPath Support](#) documentation.

For more information about Go templates and syntax, see the [Go Templates](#) documentation.