

Guided Exercise: Create Linux Containers and Kubernetes Pods

Run a base OS container in a pod and compare the environment inside the container with its host node.

Outcomes

- Create a pod with a single container, and identify the pod and its container within the container engine of an OpenShift node.
- View the logs of a running container.
- Retrieve information inside a container, such as the operating system (OS) release and running processes.
- Identify the process ID (PID) and namespaces for a container.
- Identify the User ID (UID) and supplemental group ID (GID) ranges of a project.
- Compare the namespaces of containers in one pod versus in another pod.
- Inspect a pod with multiple containers, and identify the purpose of each container.

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that all resources are available for this exercise.

```
[student@workstation ~]$ lab start pods-containers
```

Instructions

1. Log in to the OpenShift cluster and create the `pods-containers` project. Determine the UID and GID ranges for pods in the `pods-containers` project.

Log in to the OpenShift cluster as the developer user with the `oc` command.

```
[student@workstation ~]$ oc login -u developer -p developer \
https://api.ocp4.example.com:6443
Login successful
...output omitted...
```

Create the `pods-containers` project.

```
[student@workstation ~]$ oc new-project pods-containers
Now using project "pods-containers" on server "https://api.ocp4.example.com:6443".
...output omitted...
```

Identify the UID and GID ranges for pods in the `pods-containers` project.

```
[student@workstation ~]$ oc describe project pods-containers
Name:           pods-containers
Created:        16 seconds ago
Labels:         kubernetes.io/metadata.name=pods-containers
                pod-security.kubernetes.io/audit=restricted
                pod-security.kubernetes.io/audit-version=latest
                pod-security.kubernetes.io/warn=restricted
                pod-security.kubernetes.io/warn-version=latest
Annotations:   openshift.io/description=
                openshift.io/display-name=
                openshift.io/requester=developer
                openshift.io/sa.scc.mcs=s0:c28,c2
                openshift.io/sa.scc.supplemental-groups=1000760000/10000
                openshift.io/sa.scc.uid-range=1000760000/10000
Display Name:  <none>
Description:   <none>
Status:        Active
Node Selector: <none>
Quota:         <none>
Resource limits: <none>
```

Your UID and GID range values might differ from the previous output.

2. As the developer user, create a pod called `ubi9-user` from a UBI9 base container image. The image is available in the `registry.ocp4.example.com:8443/ubi9/ubi` container registry. Set the restart policy to Never and start an interactive session. Configure the pod to execute the `whoami` and `id` commands to determine the UIDs, supplemental groups, and GIDs of the container user in the pod. Delete the pod afterward.

After the ubi-user pod is deleted, log in as the admin user and then re-create the ubi9-user pod. Retrieve the UIDs and GIDs of the container user. Compare the values to the values of the ubi9-user pod that the developer user created.

Afterward, delete the ubi9-user pod.

Use the `oc run` command to create the ubi9-user pod. Configure the pod to execute the `whoami` and `id` commands through an interactive bash shell session.

```
[student@workstation ~]$ oc run -it ubi9-user --restart 'Never' \
--image registry.ocp4.example.com:8443/ubi9/ubi \
-- /bin/bash -c "whoami && id"
1000760000
uid=1000760000(1000760000) gid=0(root) groups=0(root),1000760000
```

Your values might differ from the previous output.

Notice that the user in the container has the same UID that is identified in the pods-containers project. However, the GID of the user in the container is 0, which means that the user belongs to the root group. Any files and directories that the container processes might write to must have read and write permissions by `GID=0` and have the root group as the owner.

Although the user in the container belongs to the root group, a UID value over 1000 means that the user is an unprivileged account. When a regular OpenShift user, such as the developer user, creates a pod, the containers within the pod run as unprivileged accounts.

Delete the pod.

```
[student@workstation ~]$ oc delete pod ubi9-user
pod "ubi9-user" deleted
```

Log in as the admin user with the `redhatocp` password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
Login successful.

You have access to 71 projects, the list has been suppressed. You can list all projects with 'oc projects'

Using project "pods-containers".
```

Re-create the ubi9-user pod as the admin user. Configure the pod to execute the `whoami` and `id` commands through an interactive bash shell session. Compare the values of the UID and GID for the container user to the values of the ubi9-user pod that the developer user created.

NOTE

It is safe to ignore pod security warnings when using a cluster-admin user that creates unmanaged pods. The admin user can create privileged pods that the Security Context Constraints controller does not manage.

```
[student@workstation ~]$ oc run -it ubi9-user --restart 'Never' \
--image registry.ocp4.example.com:8443/ubi9/ubi \
-- /bin/bash -c "whoami && id"
Warning: would violate PodSecurity "restricted:v1.24": allowPrivilegeEscalation != false (container "ubi9-user" must set securityContext.allowPrivilegeEscalation=false), unrestricted capabilities (container "ubi9-user" must set securityContext.capabilities.drop=["ALL"]), runAsNonRoot != true (pod or container "ubi9-user" must set securityContext.runAsNonRoot=true), seccompProfile (pod or container "ubi9-user" must set securityContext.seccompProfile.type to "RuntimeDefault" or "Localhost")
root
uid=0(root) gid=0(root) groups=0(root)
```

Notice that the value of the UID is 0, which differs from the UID range value of the pods-containers project. The user in the container is the privileged account root user and belongs to the root group. When a cluster administrator creates a pod, the containers within the pod run as a privileged account by default.

Delete the ubi9-user pod.

```
[student@workstation ~]$ oc delete pod ubi9-user
pod "ubi9-user" deleted
```

3. As the developer user, use the `oc run` command to create a ubi9-date pod from a UBI9 base container image. The image is available in the `registry.ocp4.example.com:8443/ubi9/ubi` container registry. Set the restart policy to Never, and configure the pod to execute the date command. Retrieve the logs of the ubi9-date pod to confirm that the date command executed. Delete the pod afterward.

Log in as the developer user with the developer password.

```
[student@workstation ~]$ oc login -u developer -p developer
Login successful.

You have one project on this server: "pods-containers"

Using project "pods-containers".
```

Create a pod called ubi9-date that executes the date command.

```
[student@workstation ~]$ oc run ubi9-date --restart 'Never' \
--image registry.ocp4.example.com:8443/ubi9/ubi -- date
pod/ubi9-date created
```

Wait a few moments for the creation of the pod. Then, retrieve the logs of the ubi9-date pod.

```
[student@workstation ~]$ oc logs ubi9-date
Mon Nov 28 15:02:55 UTC 2022
```

Delete the ubi9-date pod.

```
[student@workstation ~]$ oc delete pod ubi9-date
pod "ubi9-date" deleted
```

4. Use the `oc run ubi9-command -it` command to create a ubi9-command pod with the `registry.ocp4.example.com:8443/ubi9/ubi` container image. Add the `/bin/bash` in the `oc run` command to start an interactive shell. Exit the pods and view the logs for the ubi9-command pod with the `oc logs` command. Then, connect to the ubi9-command pod with the `oc attach` command, and issue the following command:

```
while true; do echo $(date); sleep 2; done
```

This command executes the date and sleep commands to generate output to the console every two seconds. Use the `oc logs` command to retrieve the logs of the ubi9 pod, and confirm that the logs display the executed date and sleep commands.

Create a pod called ubi9-command and start an interactive shell.

```
[student@workstation ~]$ oc run ubi9-command -it \
--image registry.ocp4.example.com:8443/ubi9/ubi -- /bin/bash
If you don't see a command prompt, try pressing enter.
bash-5.1$
```

Exit the shell session.

```
bash-5.1$ exit
exit
Session ended, resume using 'oc attach ubi9-command -c ubi9-command -i -t' command when the pod is running
```

Use the `oc logs` command to view the logs of the ubi9-command pod.

```
[student@workstation ~]$ oc logs ubi9-command
bash-5.1$ [student@workstation ~]$
```

The pod's command prompt is returned. The `oc logs` command displays the pod's current `stdout` and `stderr` output in the console. Because you disconnected from the interactive session, the pod's current `stdout` is the command prompt, and not the commands that you executed previously.

Use the `oc attach` command to connect to the ubi9-command pod again. In the shell, execute the `while true; do echo $(date); sleep 2; done` command to continuously generate `stdout` output.

```
[student@workstation ~]$ oc attach ubi9-command -it
If you don't see a command prompt, try pressing enter.
```

```
bash-5.1$ while true; do echo $(date); sleep 2; done
Mon Nov 28 15:15:16 UTC 2022
Mon Nov 28 15:15:18 UTC 2022
Mon Nov 28 15:15:20 UTC 2022
Mon Nov 28 15:15:22 UTC 2022
...output omitted...
```

Open another terminal window and view the logs for the ubi9-command pod with the `oc logs` command. Limit the log output to the last 10 entries with the `--tail=10` option. Confirm that the logs display the results of the command that you executed in the container.

```
[student@workstation ~]$ oc logs ubi9-command --tail=10
Mon Nov 28 15:15:16 UTC 2022
Mon Nov 28 15:15:18 UTC 2022
Mon Nov 28 15:15:20 UTC 2022
Mon Nov 28 15:15:22 UTC 2022
Mon Nov 28 15:15:24 UTC 2022
Mon Nov 28 15:15:26 UTC 2022
Mon Nov 28 15:15:28 UTC 2022
Mon Nov 28 15:15:30 UTC 2022
Mon Nov 28 15:15:32 UTC 2022
Mon Nov 28 15:15:34 UTC 2022
```

- Identify the name for the container in the ubi9-command pod. Identify the process ID (PID) for the container in the ubi9-command pod by using a debug pod for the pod's host node. Use the `crlctl` command to identify the PID of the container in the ubi9-command pod. Then, retrieve the PID of the container in the debug pod.

Identify the container name in the ubi9-command pod with the `oc get` command. Specify the JSON format for the command output. Parse the JSON output with the `jq` command to retrieve the value of the `.status.containerStatuses[].name` object.

```
[student@workstation ~]$ oc get pod ubi9-command -o json | \
  jq .status.containerStatuses[].name
"ubi9-command"
```

The ubi9-command pod has a single container of the same name.

Find the host node for the ubi9-command pod. Start a debug pod for the host with the `oc debug` command.

```
[student@workstation ~]$ oc get pods ubi9-command -o wide
NAME        READY STATUS    RESTARTS   AGE      IP          NODE      ...
ubi9-command 1/1   Running  2 (16m ago) 27m  10.8.0.26 master01 ...
```

```
[student@workstation ~]$ oc debug node/master01
Error from server (Forbidden): nodes "master01" is forbidden: User "developer" cannot get resource "nodes" in API group "" at the cluster scope
```

The debug pod fails because the developer user does not have the required permission to debug a host node.

Log in as the `admin` user with the `redhatocp` password. Start a debug pod for the host with the `oc debug` command. After connecting to the debug pod, run the `chroot /host` command to use host binaries, such as the `crlctl` command-line tool.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
Login successful.
...output omitted...
```

```
[student@workstation ~]$ oc debug node/master01
Starting pod/master01-debug ...
To use host binaries, run `chroot /host`
Pod IP: 192.168.50.10
If you don't see a command prompt, try pressing enter
```

```
sh-4.4# chroot /host
```

Use the `crlctl ps` command to retrieve the ubi9-command container ID. Specify the ubi9-command container with the `--name` option and use the JSON output format. Parse the JSON output with the `jq -r` command to get the RAW JSON output. Export the container ID as the `$CID` environment variable.

NOTE

When using `jq` without the `-r` flag, the container ID is wrapped in double quotes, which does not work with `crlctl` commands. If the `-r` flag is not used, then you can add `| tr -d '"'` to the end of the

command to trim the double quotes.

```
sh-5.1# crictl ps --name ubi9-command -o json | jq -r .containers[0].id
81adbc6222d79ed9ba195af4e9d36309c18bb71bc04b2e8b5612be632220e0d6
```

```
sh-5.1# CID=$(crictl ps --name ubi9-command -o json | jq -r .containers[0].id)
```

```
sh-5.1# echo $CID
81adbc6222d79ed9ba195af4e9d36309c18bb71bc04b2e8b5612be632220e0d6
```

Your container ID value might differ from the previous output.

Use the `crictl inspect` command to find the PID of the `ubi9-command` container. The PID value is in the `.info.pid` object in the `crictl inspect` output. Export the `ubi9-command` container PID as the `$PID` environment variable.

```
sh-5.1# crictl inspect $CID | grep pid
  "pid": 365297,
  "pids": {
    "type": "pid"
...output omitted...
  }
...output omitted...
```

```
sh-5.1# PID=365297
```

Your PID values might differ from the previous output.

6. Use the `lsns` command to list the system namespaces of the `ubi9-command` container. Confirm that the running processes in the container are isolated to different system namespaces.

View the system namespaces of the `ubi9-command` container with the `lsns` command. Specify the PID with the `-p` option and use the `$PID` environment variable. In the resulting table, the `NS` column contains the namespace values for the container.

```
sh-5.1# lsns -p $PID
      NS TYPE    NPROCS   PID USER        COMMAND
4026531835 cgroup    540     1 root      /usr/lib/systemd/systemd --swit ...
4026531837 user      540     1 root      /usr/lib/systemd/systemd --swit ...
4026536117 uts       1 153168 1000800000 /bin/bash
4026536118 ipc       1 153168 1000800000 /bin/bash
4026536120 net       1 153168 1000800000 /bin/bash
4026537680 mnt      1 153168 1000800000 /bin/bash
4026537823 pid      1 153168 1000800000 /bin/bash
```

Your namespace values might differ from the previous output.

7. Use the host debug pod to retrieve and compare the operating system (OS) and the GCC support library (`libgcc`) package version of the `ubi9-command` container and the host node.

Retrieve the OS for the host node with the `cat /etc/redhat-release` command.

```
sh-5.1# cat /etc/redhat-release
Red Hat Enterprise Linux CoreOS release 4.18
```

Use the `crictl exec` command and the `$CID` container ID variable to retrieve the OS of the `ubi9-command` container. Use the `-it` options to create an interactive terminal to execute the `cat /etc/redhat-release` command.

```
sh-5.1# crictl exec -it $CID cat /etc/redhat-release
Red Hat Enterprise Linux release 9.6 (Plow)
```

The `ubi9-command` container has a different OS from the host node.

Use the `rpm -qi libgcc` command to retrieve the `libgcc` package version of the host node.

```
sh-5.1$ rpm -qi libgcc
Name: libgcc
Version: 11.4.1
...output omitted...
```

Use the `cricctl exec` command and the `$CID` container ID variable to retrieve the `glibc` package version of the `ubi9-command` container. Use the `-it` options to create an interactive terminal to execute the `rpm -qi libgcc` command.

```
sh-5.1# crictl exec -it $CID rpm -qi libgcc
Name: libgcc
Version: 11.5.0
...output omitted...
```

The `ubi9-command` container might have a different version of the `libgcc` package from its host.

8. Exit the `master01-debug` pod and the `ubi9-command` pod.

Exit the `master01-debug` pod. You must issue the `exit` command to end the host binary access. Execute the `exit` command again to exit and remove the `master01-debug` pod.

```
sh-5.1# exit
exit

sh-4.4# exit
exit

Removing debug pod ...
Temporary namespace openshift-debug-bg7kn was removed.
```

Return to the terminal window that is connected to the `ubi9-command` pod. Press **Ctrl+C** and then execute the `exit` command. Confirm that the pod is still running.

```
...output omitted...
^C
bash-5.1$ exit
exit
Session ended, resume using 'oc attach ubi9-command -c ubi9-command -i -t' command when the pod is running
```

```
[student@workstation ~]$ oc get pods
NAME        READY   STATUS    RESTARTS   AGE
ubi9-command 1/1     Running   2 (6s ago)  35m
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish pods-containers
```