

# Chapter 4. Deploy Managed and Networked Applications on Kubernetes

[Deploy Applications from an Image and from a Template](#)

[Guided Exercise: Deploy Applications from an Image and from a Template](#)

[Manage Long-lived and Short-lived Applications by Using the Kubernetes Workload API](#)

[Guided Exercise: Manage Long-lived and Short-lived Applications by Using the Kubernetes Workload API](#)

[Kubernetes Pod and Service Networks](#)

[Guided Exercise: Kubernetes Pod and Service Networks](#)

[Scale and Expose Applications to External Access](#)

[Guided Exercise: Scale and Expose Applications to External Access](#)

[Lab: Deploy Managed and Networked Applications on Kubernetes](#)

[Summary](#)

## Abstract

|                 |   |
|-----------------|---|
| <b>Goal</b>     | Deploy applications and expose them to network access from inside and outside a Kubernetes cluster.   |
| <b>Sections</b> | <ul style="list-style-type: none"> <li>• Deploy Applications from an Image and from a Template (and Guided Exercise)</li> <li>• Manage Long-lived and Short-lived Applications by Using the Kubernetes Workload API (and Guided Exercise)</li> <li>• Kubernetes Pod and Service Networks (and Guided Exercise)</li> <li>• Scale and Expose Applications to External Access (and Guided Exercise)</li> </ul> |
| <b>Lab</b>      | <ul style="list-style-type: none"> <li>• Deploy Managed and Networked Applications on Kubernetes</li> </ul>   |

## Deploy Applications from an Image and from a Template

### Objectives

- Identify the main resources and settings that Kubernetes uses to manage long-lived applications and demonstrate how OpenShift simplifies common application deployment workflows.

## Deploying Applications

Microservices and DevOps are growing trends in enterprise software. Containers and Kubernetes gained popularity alongside those trends, and became categories of their own. Container-based infrastructures support most types of traditional and modern applications.

The *application* term can refer to a software system or to a service within it. Given this ambiguity, it is clearer to refer directly to resources, services, and other components.

## Resources and Resource Definitions

Kubernetes manages applications, or services, as a loose collection of resources. Resources are configuration pieces for components in the cluster. When you create a resource, the corresponding component is not created immediately. Instead, the cluster is responsible for eventually creating the component.

A *resource type* represents a specific component type, such as a pod. Kubernetes ships with many default resource types, some of which overlap in function. Red Hat OpenShift Container Platform (RHOC) includes the default Kubernetes resource types, and provides other resource types of its own. To add resource types, you can create or import CRDs.

## Managing Resources

You can add, view, and edit resources in various formats, including YAML and JSON format. Traditionally, YAML is the most common format.

You can delete resources in batch by using label selectors or by deleting the entire project or namespace. For example, the following command deletes only deployments with the app=my-app label:

```
[user@host ~]$ oc delete deployment -l app=my-app
deployment.apps "my-app" deleted
```

Similar to creation, deleting a resource is not immediate, but is instead a request for eventual deletion.

**NOTE**

Commands that are executed without specifying a namespace are executed in the user's current namespace.

## Common Resource Types and Their Uses

The following resources are standard OpenShift resources:

### Templates

Similar to projects, templates are an RHOCP addition to Kubernetes. A template is a YAML manifest that contains parameterized definitions of one or more resources. RHOCP provides predefined templates in the openshift namespace.

You can process a template into a list of resources by using the `oc process` command, which replaces values and generates resource definitions. The resulting resource definitions create or update resources in the cluster by supplying them to the `oc apply` command.

For example, the following command processes the `mysql-template.yaml` template file and generates four resource definitions:

```
[user@host ~]$ oc process -f mysql-template.yaml -o yaml
apiVersion: v1
items:
- apiVersion: v1
  kind: Secret
  ...output omitted...
- apiVersion: v1
  kind: Service
  ...output omitted...
- apiVersion: v1
  kind: PersistentVolumeClaim
  ...output omitted...
- apiVersion: apps/v1
  kind: Deployment
  ...output omitted...
```

The `--parameters` option instead displays the parameters of a template. For example, the following command lists the parameters of the `mysql-template.yaml` file.

```
[user@host ~]$ oc process -f mysql-template.yaml --parameters
NAME          DESCRIPTION ...output omitted...
...output omitted...
MYSQL_USER     Username for MySQL user    ...output omitted...
MYSQL_PASSWORD Password for the MySQL connection ...output omitted...
MYSQL_ROOT_PASSWORD Password for the MySQL root user. ...output omitted...
MYSQL_DATABASE Name of the MySQL database accessed. ...output omitted...
VOLUME_CAPACITY Volume space available for data,    ...output omitted...
```

You can use templates with the `new-app` command from RHOCP. In the following example, the `new-app` command uses the `mysql-persistent` template to create a MySQL application and its supporting resources:

```
[user@host ~]$ oc new-app --template mysql-persistent
--> Deploying template "db-app/mysql-persistent" to project db-app
...output omitted...
The following service(s) have been created in your project: mysql.

  Username: userQSL
  Password: pyf0yElPvFWYQQou
  Database Name: sampledb
  Connection URL: mysql://mysql:3306/
...output omitted...
* With parameters:
  * Memory Limit=512Mi
  * Namespace=openshift
  * Database Service Name=mysql
  * MySQL Connection Username=userQSL # generated
  * MySQL Connection Password=pyf0yElPvFWYQQou # generated
  * MySQL root user Password=HHbdurqW05gAog2m # generated
  * MySQL Database Name=sampledbs
  * Volume Capacity=1Gi
  * Version of MySQL Image=8.0-e18

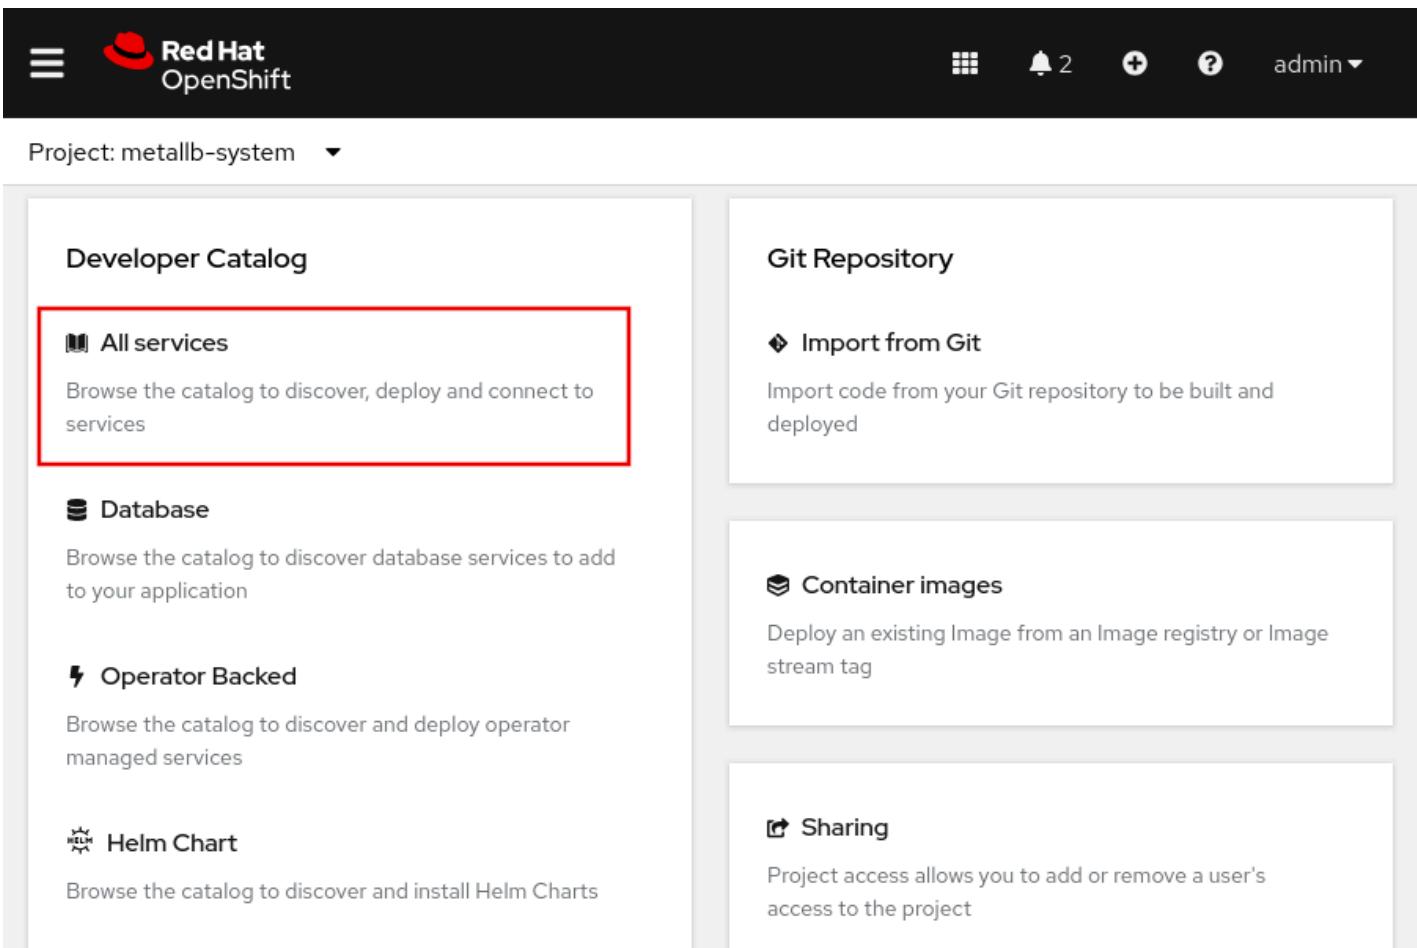
--> Creating resources ... ⓘ
secret "mysql" created
service "mysql" created
persistentvolumeclaim "mysql" created
deployment.apps "mysql" created
--> Success
...output omitted...
```

- 1 Several resources are created to meet the requirements of the deployment, including a secret, a service, and a persistent volume claim.

#### NOTE

You can specify environment variables to be configured in creating your application.

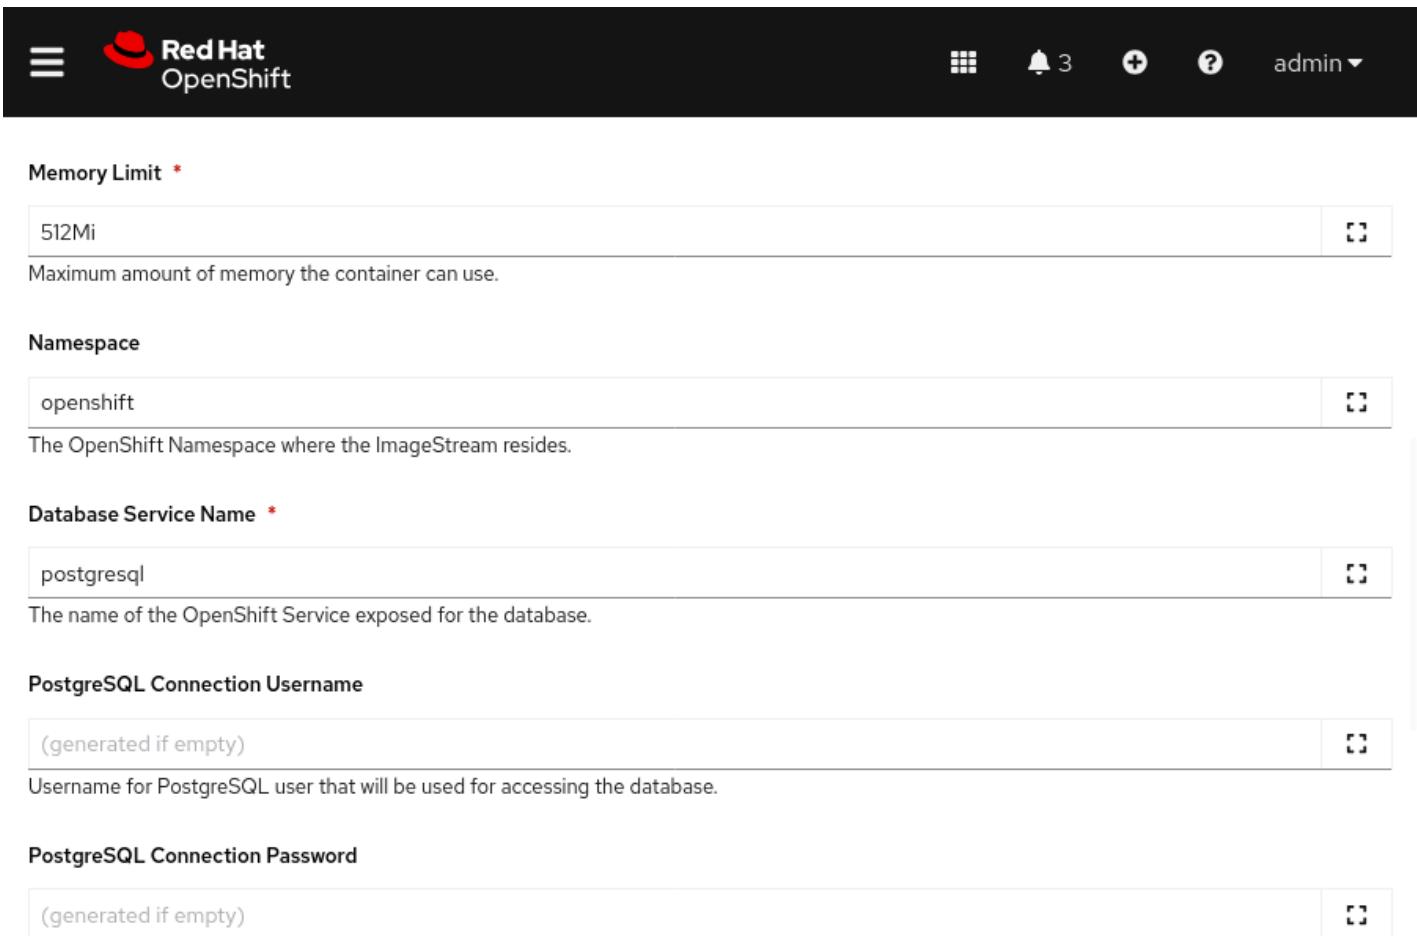
You can also use templates, helm charts, builder images, or devfiles to create applications from the web console by using the **Developer Catalog** page. From the **Developer** perspective, go to the **+Add** menu and click **All Services** in the **Developer Catalog** section to open the catalog.



The screenshot shows the Red Hat OpenShift Developer Catalog interface. On the left, there's a sidebar with a 'Project: metallb-system' dropdown. The main area has several sections:

- Developer Catalog**: Contains a red-bordered box around the **All services** section.
  - All services**: Description: "Browse the catalog to discover, deploy and connect to services".
  - Database**: Description: "Browse the catalog to discover database services to add to your application".
  - Operator Backed**: Description: "Browse the catalog to discover and deploy operator managed services".
  - Helm Chart**: Description: "Browse the catalog to discover and install Helm Charts".
- Git Repository**: Contains a red-bordered box around the **Import from Git** section.
  - Import from Git**: Description: "Import code from your Git repository to be built and deployed".
  - Container images**: Description: "Deploy an existing Image from an Image registry or Image stream tag".
  - Sharing**: Description: "Project access allows you to add or remove a user's access to the project".

Then, enter the application name in the filter box to search for an application's template. You can instantiate a template and change the default values, such as the Git repository, the memory limits, and the application version.



The screenshot shows the configuration page for a PostgreSQL application. At the top, there's a navigation bar with the Red Hat OpenShift logo and a 'Memory Limit \*' field set to '512Mi'. Below it are sections for 'Namespace' (set to 'openshift') and 'Database Service Name \*' (set to 'postgresql'). Further down are fields for 'PostgreSQL Connection Username' (placeholder '(generated if empty)') and 'PostgreSQL Connection Password' (placeholder '(generated if empty)').

You can also create an application from the Topology menu, and by clicking **Start building application** or clicking the book icon.

The screenshot shows the Red Hat OpenShift web interface. At the top, there is a navigation bar with the Red Hat logo, the text "Red Hat OpenShift", and user information for "admin". Below the navigation bar, the main content area has a header with "Project: pods-review" and "Application: All applications". There are several search and filter options: "Display options", "Filter by resource", "Name", and "Find by name...". A red box highlights the "Display options" button. In the center, there is a large, stylized molecular or network graph icon. Below it, the text "No resources found" is displayed. A red box highlights the blue link "Start building your application" which is followed by the text "or visit the [Add page](#) for more details." At the bottom of the screen, there are four small icons: a magnifying glass, a plus sign, a minus sign, and a refresh symbol.

## Pod

A pod is the smallest compute unit that you can define, deploy, and manage in RHOCP. A pod runs one or more containers that represent a single application. Containers in the pod share resources, such as networking and storage.

The following example shows the definition of a pod:

```

apiVersion: v1
kind: Pod 1
metadata: 2
  annotations: {}
  labels:
    deployment: docker-registry-1
  name: registry
  namespace: pod-registries
spec: 3
  containers:
    - env:
      - name: OPENSHIFT_CA_DATA
        value:
      image: openshift/origin-docker-registry:v0.6.2
      imagePullPolicy: IfNotPresent
      name: registry
      ports:
        - containerPort: 5000
          protocol: TCP
      resources: {}
      securityContext: {}
      volumeMounts:
        - mountPath: /registry
          name: registry-storage
    dnsPolicy: ClusterFirst
    imagePullSecrets:
      - name: default-dockercfg-at06w
    restartPolicy: Always
    serviceAccount: default
    volumes:
      - emptyDir: {}
        name: registry-storage
  status: 4
    conditions: {}

```

- 1** The resource kind is set to Pod.
- 2** Information that describes your application, such as the name, project, attached labels, and annotations.
- 3** Section where the application requirements are specified, such as the container name, the container image, environment variables, volume mounts, network configuration, and volumes.
- 4** Indicates the last condition of the pod, such as the last probe time, the last transition time, the status setting as true or false, and more.

## Deployments

A deployment describes the intended state of a component of the application as a pod template. A deployment manages one or more replica sets.

The following example shows the definition of a deployment:

```

apiVersion: apps/v1
kind: Deployment 1
metadata:
  name: hello-openshift 2
spec:
  replicas: 1 3
  selector:
    matchLabels:
      app: hello-openshift
  template: 4
    metadata:
      labels:
        app: hello-openshift
    spec:
      containers:
        - name: hello-openshift 5
          image: openshift/hello-openshift:latest 6
          ports: 7
            - containerPort: 80

```

- 1** The resource kind is set to Deployment.
- 2** Name of the deployment resource.

- ③ Number of running instances.
- ④ Section to define the metadata, labels, and the container information of the deployment resource.
- ⑤ Name of the container.
- ⑥ Resource of the image for creating the deployment resource.
- ⑦ Port configuration, such as the port number, the name of the port, and the protocol.

## Projects

RHOCP adds projects to enhance the function of Kubernetes namespaces. A project is a Kubernetes namespace with additional annotations, and is the primary method for managing access to resources for regular users. Projects can be created from templates and must use RBAC for organization and permission management. Administrators must grant cluster users access to a project. If a cluster user is allowed to create projects, then the user automatically has access to their created projects.

Projects provide logical and organizational isolation to separate your application component resources. Resources in one project can access resources in other projects, but not by default.

The following example shows the definition of a project:

```
apiVersion: project.openshift.io/v1
kind: Project ①
metadata:
  name: test ②
spec:
  finalizers: ③
    - kubernetes
```

- ① The resource kind is set to Project.
- ② Name of the project.
- ③ A finalizer is a special metadata key that tells Kubernetes to wait until a specific condition is met before it fully deletes a resource.

## Services

You can configure internal pod-to-pod network communication in RHOCP by using a service. Applications send requests to the service name and port. RHOCP provides a virtual network, which reroutes such requests to the pods that the service targets by using labels.

The following example shows the definition of a service:

```
apiVersion: v1
kind: Service ①
metadata:
  name: docker-registry ②
  namespace: test ③
spec:
  selector:
    app: MyApp ④
  ports:
    - protocol: TCP ⑤
      port: 80 ⑥
      targetPort: 9376 ⑦
```

- ① The resource kind is set to Service.
- ② Name of the service.
- ③ Project name where the service resource exists.
- ④ The label selector identifies all pods with the attached app=MyApp label and adds the pods to the service endpoints.
- ⑤ Internet protocol is set to TCP.
- ⑥ Port that the service listens on.
- ⑦ Port on the backing pods, which the service forwards connections to.

## Persistent Volume Claims

RHOCP uses the Kubernetes persistent volume (PV) framework to enable cluster administrators to provision persistent storage for a cluster. Developers can use persistent volume claims (PVCs) to request PV resources without having specific knowledge of the underlying storage infrastructure. After a PV is bound to a PVC, that PV cannot then be bound to additional PVCs. Because PVCs are namespaced objects and PVs are globally scoped objects, this binding effectively scopes a bound PV to a single namespace until the binding PVC is deleted.

The following example shows the definition of a persistent volume claim:

```
apiVersion: v1
kind: PersistentVolumeClaim 1
metadata:
  name: mysql-pvc 2
spec:
  accessModes:
    - ReadWriteOnce 3
  resources:
    requests:
      storage: 1Gi 4
  storageClassName: nfs-storage 5
status: {}
```

- 1** The resource kind is set to `PersistentVolumeClaim`.
- 2** Name of the service.
- 3** The access mode, to define the read/write and mount permissions.
- 4** The storage of the PVC.
- 5** Name of the StorageClass that the claim requires.

## Secrets

Secrets provide a mechanism to hold sensitive information, such as passwords, private source repository credentials, sensitive configuration files, and credentials to an external resource, such as an SSH key or an OAuth token. You can mount secrets into containers by using a volume plug-in. Kubernetes can also use secrets to perform actions, such as on pods, declaring environment variables. Secrets can store any type of data. Kubernetes and OpenShift support different types of secrets, such as service account tokens, SSH keys, and TLS certificates.

The following example shows the definition of a secret:

```
apiVersion: v1
kind: Secret 1
metadata:
  name: example-secret 2
  namespace: my-app 3
type: Opaque 4
data: 5
  username: bXl1c2Vycg==
  password: bXLQQDU1cg==
stringData: 6
  hostname: myapp.mydomain.com
  secret.properties:
    property1=valueA
    property2=valueB
```

- 1** The resource kind is set to `Secret`.
- 2** Name of the service.
- 3** Project name where the service resource exists.
- 4** Specifies the type of secret.
- 5** Specifies the encoded string and data.
- 6** Specifies the decoded string and data.

## Managing Resources from the Command Line

Many Kubernetes and RHOCP commands can create and modify cluster resources. Some commands are part of core Kubernetes, whereas others are exclusive additions to RHOCP.

The resource management commands generally fall into one of two categories:

- a. An *imperative* command instructs what the cluster does.
- b. A *declarative* command defines the state that the cluster attempts to match.

## Imperative Resource Management

The `create` command is an imperative way to create resources, and is included in both of the `oc` and `kubectl` commands. For example, the following command creates a deployment named `my-app` that creates pods that are based on the specified image.

```
[user@host ~]$ oc create deployment my-app --image example.com/my-image:dev
deployment.apps/my-app created
```

Use the `set` command to define attributes on a resource, such as environment variables. For example, the following command adds the `TEAM=red` environment variable to the preceding deployment.

```
[user@host ~]$ oc set env deployment/my-app TEAM=red
deployment.apps/my-app updated
```

Another imperative approach to creating a resource is the `run` command. In the following example, the `run` command creates the `example-pod` pod.

```
[user@host ~]$ oc run example-pod \
--image=registry.access.redhat.com/ubi8/httpd-24 \
--env GREETING='Hello from the awesome container' \
--port 8080
pod/example-pod created
```

- ➊ The pod `.metadata.name` definition
- ➋ The image for the single container in this pod
- ➌ The environment variable for the single container in this pod
- ➍ The port metadata definition

The imperative commands are a faster way of creating pods, because such commands do not require a pod object definition. Versioning and incremental changes require declarative manifests.

Generally, developers test a deployment by using imperative commands, and then use the imperative commands to generate the pod object definition. Use the `--dry-run=client` option to avoid creating the object in RHOCP. Additionally, use the `-o yaml` or `-o json` option to configure the definition format.

The following command is an example of generating the YAML definition for the `example-pod` pod:

```
[user@host ~]$ oc run example-pod \
--image=registry.access.redhat.com/ubi8/httpd-24 \
--env GREETING='Hello from the awesome container' \
--port 8080 \
--dry-run=client -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: example-pod
    name: example-pod
spec:
  containers:
    ...output omitted...
```

Managing resources in this way is imperative, because you are instructing the cluster what to do rather than declaring the intended outcomes.

## Declarative Resource Management

Using the `create` command does not guarantee that you are creating resources imperatively. For example, providing a manifest declaratively creates the resources that are defined in the YAML file, such as in the following command.

```
[user@host ~]$ oc create -f my-app-deployment.yaml
deployment.apps/my-app created
```

RHOCP adds the new-app command, which provides another declarative way to create resources. This command uses heuristics to automatically determine which types of resources to create based on the specified parameters. For example, the following command deploys the my-app application by creating several resources, from a YAML manifest file:

```
[user@host ~]$ oc new-app --file=./example/my-app.yaml
...output omitted...
--> Creating resources ...
imagestream.image.openshift.io "my-app" created
deployment.apps "my-app" created
services "my-app" created
...output omitted...
```

In both of the preceding create and new-app examples, you are declaring the intended state of the resources, and so they are declarative.

You can also use the new-app command with templates for resource management. A template describes the intended state of resources that must be created for an application to run, such as deployments and services. Supplying a template to the new-app command is a form of declarative resource management.

The new-app command also includes options, such as the --param option, that customize an application deployment declaratively. For example, when the new-app is used with a template, you can include the --param option to override a parameter value in the template.

```
[user@host ~]$ oc new-app --template mysql-persistent \
--param MYSQL_USER=operator --param MYSQL_PASSWORD=myP@55 \
--param MYSQL_DATABASE=mydata \
--param DATABASE_SERVICE_NAME=db
--> Deploying template "db-app/mysql-persistent" to project db-app
...output omitted...
The following service(s) have been created in your project: db.

      Username: operator
      Password: myP@55
      Database Name: mydata
      Connection URL: mysql://db:3306/
...output omitted...
* With parameters:
  * Memory Limit=512Mi
  * Namespace=openshift
  * Database Service Name=db
  * MySQL Connection Username=operator
  * MySQL Connection Password=myP@55
  * MySQL root user Password=tLh8BThuVgnIrCon # generated
  * MySQL Database Name=mydata
  * Volume Capacity=1Gi
  * Version of MySQL Image=8.0-1e8

--> Creating resources ...
secret "db" created
service "db" created
persistentvolumeclaim "db" created
--> Success
...output omitted...
```

Similar to the create command, you can use the new-app command imperatively. When using a container image with the new-app command, you are instructing the cluster what to do, rather than declaring the intended outcomes. For example, the following command deploys the example.com/my-app:dev image by creating several resources, including a deployment resource:

```
[user@host ~]$ oc new-app --image example.com/my-app:dev
...output omitted...
--> Creating resources ...
imagestream.image.openshift.io "my-app" created
deployment.apps "my-app" created
...output omitted...
```

You can also supply a Git repository to the new-app command. The following command creates an application named httpd24 by using a Git repository:

```
[user@host ~]$ oc new-app https://github.com/apache/httpd.git#2.4.56
...output omitted...
--> Creating resources ...
imagestream.image.openshift.io "httpd24" created
deployment.apps "httpd24" created
...output omitted...
```

## Retrieving Resource Information

You can supply the `all` argument to commands, to specify a list of common resource types. However, the `all` option does not include every resource type. Instead, `all` is a shorthand form for a predefined subset of types. When you use this command argument, ensure that `all` includes any intended types to address.

You can view detailed information about a resource, such as the defined parameters, by using the `describe` command. For example, RHOCP provides templates in the `openshift` project to use with the `oc new-app` command. The following example command displays detailed information about the `mysql-ephemeral` template:

```
[user@host ~]$ oc describe template mysql-ephemeral -n openshift
Name:          mysql-ephemeral
Namespace:     openshift
...output omitted...
Parameters:
  Name:           MEMORY_LIMIT
  Display Name:  Memory Limit
  Description:   Maximum amount of memory the container can use.
  Required:      true
  Value:         512Mi

  Name:           NAMESPACE
  Display Name:  Namespace
  Description:   The OpenShift Namespace where the ImageStream resides.
  Required:      false
  Value:         openshift
...output omitted...
Objects:
  Secret          ${DATABASE_SERVICE_NAME}
  Service         ${DATABASE_SERVICE_NAME}
```

The `oc describe` command cannot generate structured output, such as the YAML or JSON formats. A structured format is required for the `oc describe` command to parse or filter the output with tools such as JSONPath or Go templates." Instead, use the `oc get` command to generate and to parse the structured output of a resource.

### REFERENCES

- [OpenShift Container Platform Documentation - Understanding Deployments](#)
- [OpenShift Container Platform Documentation - Working with Projects](#)
- [OpenShift Container Platform Documentation - Creating Applications Using the CLI](#)
- [OpenShift Container Platform Documentation - Using Pods](#)