

Scale and Expose Applications to External Access

Objectives

- Expose applications to clients outside the cluster by using Kubernetes ingress and OpenShift routes.

Outcomes

- Allocate IP addresses for pods and services in a Red Hat OpenShift Container Platform (RHOCP) cluster.
- Configure services to enable communication between pods.
- Expose applications to external networks by using routes and ingress objects.

IP Addresses for Pods and Services

Real-world applications typically require multiple pods to scale horizontally. To meet growing user demand, RHOCP deploys many pods that run the same containers from a single pod resource definition. A Service resource defines a single IP and port combination. The resource assigns a stable IP address to a pool of pods. The resource balances client requests across those pods.

RHOCP assigns each service a unique IP address for clients to connect to. The assignment uses a round-robin load-balancing strategy by default. The service IP address originates from an internal OVN-Kubernetes network. The network, which is separate from the pod network, is accessible only to pods within the cluster. RHOCP adds each pod to match the selector of the service as an endpoint.

NOTE

OVN-Kubernetes, which is the default network provider starting with RHOCP 4.18, manages pod-to-pod and service-to-pod communication. The provider ensures efficient traffic routing within the cluster.

Containers within Kubernetes pods must avoid connecting directly to dynamic IP addresses of other pods. Instead, services link stable OVN-Kubernetes IP addresses to pods to ensure scalability and fault tolerance. When pods restart, replicate, or reschedule to different nodes, RHOCP updates the service endpoints accordingly.

Service Types

RHOCP offers several service types to meet diverse application needs, cluster infrastructure, and security requirements.

ClusterIP

The `ClusterIP` type, which is the default unless otherwise specified, assigns a cluster-internal IP address to the service. The assignment makes the service accessible only within the RHOCP cluster.

`clusterIP` services enable pod-to-pod communication. RHOCP allocates IP addresses from a dedicated service network. The network, which is accessible only inside the cluster, supports internal communication. Most applications benefit from this service type because Kubernetes automates management of the service.

LoadBalancer

The `LoadBalancer` type directs RHOCP to provision a load balancer of the cloud provider. The load balancer assigns an externally accessible IP address to the application.

Use caution when deploying `LoadBalancer` services because costs can be high when assigning a load balancer to every application. These services expose applications to external networks. The exposure requires additional security configurations to prevent unauthorized access.

ExternalIP

The `ExternalIP` service redirects traffic from a virtual IP address on a cluster node to a pod. A cluster administrator assigns the virtual IP address to a node. The administrator configures failover to another node if needed.

WARNING

Starting with RHOCP 4.18, `ExternalIP` is a legacy feature due to security concerns. Direct network connections to cluster nodes, which this service requires, conflict with most security policies. Use the `LoadBalancer` or the `NodePort` type unless the `ExternalIP` type meets specific requirements.

RHOCP configures Network Address Translation (NAT) rules to route traffic from the virtual IP address to the pod. Administrators must ensure that external IP addresses route correctly to nodes. Additional security measures protect the cluster from external access.

NodePort

The `NodePort` service exposes a service on a specific port, with a default range of 30000–32767, on the IP address of each cluster node. Each node redirects traffic from this port to the endpoints of the service.

WARNING

NodePort services require direct network connections to cluster nodes, which poses a security risk. Implement strict firewall rules and access controls to mitigate risks.

ExternalName

The ExternalName service maps a Kubernetes service to an external DNS name that is specified in the externalName field. The Kubernetes DNS server returns the externalName in a Canonical Name (CNAME) record. The record directs clients to resolve the external name to an IP address.

Using Routes for External Connectivity

RHOCP provides Route resources to expose applications to external networks. Route resources support HTTP, HTTPS, TCP, and other protocols, and prioritize HTTP and TLS-based applications for external exposure. Non-HTTP applications, such as databases, typically remain internal. Route and Ingress resources handle ingress traffic in RHOCP.

A Route resource assigns a publicly accessible hostname to an application. The OpenShift ingress controller, which is HAProxy-based, redirects traffic from a public IP address to pods. The OpenShift ingress controller is the default, with support for third-party ingress controllers that are deployed in parallel. Route resources offer advanced features, which surpass standard Kubernetes Ingress capabilities. Features include TLS re-encryption, passthrough, and split traffic for blue-green deployments.

To create a route by using the oc CLI, use the following command:

```
[user@host ~]$ oc expose service api-frontend --hostname api.apps.acme.com
```

The --target-port option of the oc expose command specifies the name or number of the container port that the service uses to communicate with the pods. If you provide no port value, then the port copies from the configuration of the deployment.

Omitting the --hostname option prompts RHOCP to generate a hostname in the <route-name>-<project-name>. <default-domain> format. For example, a frontend route in an api project with a wildcard domain of apps.example.com becomes frontend-api.apps.example.com.

IMPORTANT

The DNS server for the wildcard domain resolves all names to configured IPs. The server does not recognize specific route hostnames. The OpenShift ingress controller treats route hostnames as HTTP virtual hosts. Invalid or non-existent route hostnames trigger an HTTP 503 error.

When creating a route, configure the following settings:

1. **Service Name:** Specifies the service to determine target pods.
2. **Hostname:** Sets the route hostname as a subdomain of the cluster wildcard domain, which RHOCP can autogenerate.
3. **Path:** Defines an optional path for path-based routing.
4. **Target Port:** Matches the targetPort in the service, and specifies the listening port of the application.
5. **Encryption Strategy:** Chooses secure (TLS) or insecure routing.

The following example shows a minimal route definition:

```
kind: Route
apiVersion: route.openshift.io/v1
metadata:
  name: a-simple-route 1
  labels: 2
    app: API
    name: api-frontend
spec:
  host: api.apps.acme.com 3
  to:
    kind: Service
    name: api-frontend 4
    port:
      targetPort: 8443 5
```

- 1** Specifies a unique route name.
- 2** Defines selectors for the route.
- 3** Sets the hostname, which is a subdomain of the cluster wildcard domain.
- 4** Identifies the service to redirect traffic to, by determining target pods.
- 5** Maps the router port to the endpoint port of the service.

NOTE

Some ecosystem components integrate with Ingress resources but not with Route resources. Starting with RHOCP 4.18, creating an Ingress object automatically generates a managed Route object, which is deleted when the Ingress object is removed.

To delete a route, use the following command:

```
[user@host ~]$ oc delete route a-simple-route
```

To create a route from the RHOPC web console, go to **Networking → Routes**. Click **Create Route**. Customize the name, hostname, path, and target service by using the form or the YAML editor.

The screenshot shows the 'Create Route' interface in the OpenShift web console. The project dropdown at the top is set to 'metallb-system'. The main section is titled 'Create Route' with a sub-instruction: 'Routing is a way to make your application publicly visible.' There are two options for configuration: 'Form view' (selected) and 'YAML view'. The 'Name' field is required and has 'my-route' entered. The 'Hostname' field has 'www.example.com' entered. The 'Path' field is empty. At the bottom are two buttons: a grey 'Create' button and a blue 'Cancel' button.

Using Ingress Objects for External Connectivity

A Kubernetes Ingress resource provides similar functions to RHOPC Route resources, and supports HTTP, HTTPS, Server Name Indication (SNI), and TLS with SNI. The OpenShift ingress controller processes Ingress objects, and enables features such as TLS termination, path redirection, and sticky sessions.

NOTE

Although Ingress resources are standard in Kubernetes, Route resources are preferred for advanced features and native integration with the OpenShift ingress controller.

To create an Ingress object, use the following command:

```
[user@host ~]$ oc create ingress ingr-sakila --rule="ingr-sakila.apps.ocp4.example.com/*=sakila-service:8080"
```

The following example shows a minimal Ingress definition:

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: frontend 1
spec:
  rules:
  - host: www.example.com 2
    http:
      paths:
      - path: /
        pathType: Prefix
      backend:
        service:
          name: frontend 3
          port:
            number: 80 4
  tls:
  - hosts:
    - www.example.com
    secretName: example-com-tls-certificate 5

```

- 1** Specifies a unique name for the Ingress object.
- 2** Sets the hostname for inbound traffic.
- 3** Identifies the service to redirect traffic to.
- 4** Defines the port for the service back end.
- 5** Configures TLS for secure paths, which requires a matching host and a certificate secret.

To delete an Ingress object, use the following command:

```
[user@host ~]$ oc delete ingress frontend
```

Sticky Sessions

Sticky sessions ensure that stateful application traffic reaches the same pod for a user session. The OpenShift ingress controller can use cookies to manage session persistence for Route and Ingress resources. The controller selects a pod for a user request. The controller generates a cookie. The controller includes the cookie in the response. The client sends the cookie with subsequent requests, which enables the controller to route traffic to the same pod.

To configure session persistence for an Ingress object by using a cookie, use the following command:

```
[user@host ~]$ oc annotate ingress ingr-example ingress.kubernetes.io/affinity=cookie
```

To configure session persistence for a Route object with a custom cookie named `myapp`, use the following command:

```
[user@host ~]$ oc annotate route route-example router.openshift.io/cookie_name=myapp
```

To capture the route hostname, use the following command:

```
[user@host ~]$ ROUTE_NAME=$(oc get route route-example -o jsonpath='{.spec.host}')
```

To access the route and save the cookie, use the following command:

```
[user@host ~]$ curl $ROUTE_NAME -k -c /tmp/cookie_jar
```

To use the saved cookie for subsequent requests, use the following command:

```
[user@host ~]$ curl $ROUTE_NAME -k -b /tmp/cookie_jar
```

Using the saved cookie ensures that requests reach the same pod.

NOTE

Cookies cannot be set on passthrough routes because the HTTP traffic is encrypted, and the router does not terminate TLS or read the contents of the request.

Load Balance and Scale Applications

Administrators and developers can scale the number of replica pods in a deployment to handle traffic surges or to conserve resources.

To scale a deployment, use the following command:

```
[user@host ~]$ oc scale --replicas=5 deployment/scale
```

The deployment updates the replica set. The replica set creates or deletes pods to match the intended replica count.

Avoid modifying a replica set directly, because the deployment controller manages these changes.

Load Balance Pods

A Kubernetes Service resource functions as an internal load balancer, by providing access to a workload via the service name. The service distributes incoming connections across replicated pods. The OpenShift ingress controller uses the selector of the service to identify the service and endpoints. When both the controller and a service provide load balancing, RHOCP relies on the controller to direct traffic to pods. The controller detects changes in IP addresses of the service. The controller updates the configuration. Custom controllers can propagate API object changes to external routing solutions.

The OpenShift ingress controller maps external hostnames. The controller balances service endpoints over protocols that include distinguishing information, such as HTTP host headers.

REFERENCES

For more information, see the *Understanding Networking* section in the Red Hat OpenShift Container Platform 4.18 *Networking* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/networking_overview/understanding-networking

For more information about services, refer to the *Service [v1]* section in the Red Hat OpenShift Container Platform 4.18 *Network APIs* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/network_apis/service-v1

For more information about routes, refer to the *Configuring Routes* section in the Red Hat OpenShift Container Platform 4.18 *Ingress and Load Balancing* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/ingress_and_load_balancing/configuring-routes

For more information about load balancing, refer to the *Ingress and Load Balancing* section in the Red Hat OpenShift Container Platform 4.18 *Ingress and Load Balancing* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/ingress_and_load_balancing/index

For more information, see the *Cluster Network Operator in OpenShift Container Platform* chapter in the Red Hat OpenShift Container Platform 4.18 *Networking* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/networking_operators/cluster-network-operator

For more information, see the *About the OVN-Kubernetes Network Plug-in* chapter in the Red Hat OpenShift Container Platform 4.18 *Networking* documentation at https://docs.redhat.com/en/documentation/openshift_container_platform/4.18/html/ovn-kubernetes_network_plugin/index