

Guided Exercise: Manage Non-shared Storage with Stateful Sets

Deploy a replicated web server by using a deployment and verify that all web server pods share a PV; and deploy a replicated MySQL database by using a stateful set and verify that each database instance gets a dedicated PV.

Outcomes

- Deploy a web server with persistent storage.
- Add data to the persistent storage.
- Scale the web server deployment and observe the data that is shared with the replicas.
- Create a database server with a headless service and a stateful set by using YAML manifest files.
- Verify that each instance from the stateful set has a persistent volume claim.

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that all resources are available for this exercise.

```
[student@workstation ~]$ lab start storage-statefulsets
```

Instructions

1. Create a web server deployment named `web-server`. Use the `registry.ocp4.example.com:8443/redhattraining/hello-world-nginx:latest` container image.

Log in to the OpenShift cluster as the developer user with developer as the password.

```
[student@workstation]$ oc login -u developer -p developer \
https://api.ocp4.example.com:6443
...output omitted...
```

Change to the `storage-statefulsets` project.

```
[student@workstation]$ oc project storage-statefulsets
Now using project "storage-statefulsets" on server ...output omitted...
```

Create the `web-server` deployment.

```
[student@workstation ~]$ oc create deployment web-server \
--image registry.ocp4.example.com:8443/redhattraining/hello-world-nginx:latest
deployment.apps/web-server created
```

Verify the deployment status.

```
[student@workstation ~]$ oc get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
web-server 1/1     1           1           50s
```

2. Add the PVC named `web-pv` to the `web-server` deployment. Use the default storage class and the following information to create the persistent volume claim:

Field	Value
Name	web-pv
Type	persistentVolumeClaim
Claim mode	rwo
Claim size	5Gi
Mount path	/usr/share/nginx/html
Claim name	web-pv-claim

Use the `oc set volumes` command to create and add the `web-pv` persistent volume to the `web-server` deployment.

```
[student@workstation ~]$ oc set volumes deployment/web-server \
--add --name web-pv --type persistentVolumeClaim --claim-mode rwo \
--claim-size 5Gi --mount-path /usr/share/nginx/html --claim-name web-pv-claim
deployment.apps/web-server volume updated
```

Because a storage class was not specified with the `--claim-class` option, the command uses the default storage class to create the persistent volume.

Verify the deployment pod status. Notice that a new pod is created and in running status.

```
[student@workstation ~]$ oc get pods -l app=web-server
NAME           READY   STATUS    RESTARTS   AGE
web-server-5fb94fd568-48n8t   1/1     Running   0          5s
```

Check that the web-server pod mounts the web-pv PVC. The pod name is different in your cluster.

```
[student@workstation ~]$ oc get pod web-server-5fb94fd568-48n8t -o json | \
jq .spec.volumes[0].persistentVolumeClaim.claimName
"web-pv-claim"
```

Verify that the web-pv PVC shows the bound status.

```
[student@workstation ~]$ oc get pvc
NAME      STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS ...
web-pv-claim Bound    pvc-42...  5Gi        RWO          nfs-storage ...
```

The default storage class, nfs-storage, provisioned the persistent volume.

3. Add data to the web-pv PVC by using the exec command.

List the pods to retrieve the web-server pod name.

```
[student@workstation ~]$ oc get pods
NAME           READY   STATUS    RESTARTS   AGE
web-server-5fb94fd568-48n8t   1/1     Running   0          3m
```

The pod name is different in your cluster.

Use the exec command to add the pod name that you retrieved from the previous step to the `/usr/share/nginx/html/index.html` file on the pod.

```
[student@workstation ~]$ oc exec -it pod/web-server-5fb94fd568-48n8t \
-- /bin/bash -c \
'echo "Hello, World from ${HOSTNAME}" > /usr/share/nginx/html/index.html'
```

Retrieve the contents of the `/var/www/html/index.html` file to confirm that the pod name is in the file.

```
[student@workstation ~]$ oc exec -it pod/web-server-5fb94fd568-48n8t \
-- cat /usr/share/nginx/html/index.html
Hello, World from web-server-5fb94fd568-48n8t
```

4. Scale the web-server deployment to two replicas and confirm that an additional pod is created.

Scale the web-server deployment to two replicas.

```
[student@workstation ~]$ oc scale deployment web-server --replicas 2
deployment.apps/web-server scaled
```

Verify the replica status and retrieve the pod names.

```
[student@workstation ~]$ oc get pods
NAME           READY   STATUS    RESTARTS   AGE
web-server-5fb94fd568-48n8t   1/1     Running   0          10m
web-server-5fb94fd568-rqvpb   1/1     Running   0          6s
```

The pod names are different in your cluster.

5. Retrieve the content of the `/usr/share/nginx/html/index.html` file on the web-server pods by using the `oc exec` command to verify that the file is the same in both pods.

Verify that the /usr/share/nginx/html/index.html file is the same in both pods.

```
[student@workstation ~]$ oc exec -it pod/web-server-5fb94fd568-48n8t \
-- cat /usr/share/nginx/html/index.html
Hello, World from web-server-5fb94fd568-48n8t
```

```
[student@workstation ~]$ oc exec -it pod/web-server-5fb94fd568-rqvpb \
-- cat /usr/share/nginx/html/index.html
Hello, World from web-server-5fb94fd568-48n8t
```

Notice that both files show the name of the first instance, because they share the persistent volume.

6. Create a database server by using a stateful set and a headless service.

Edit the /home/student/DO180/labs/storage-statefulsets/statefulset-db.yaml file and modify the placeholder text to match the following content:

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: dbserver
spec:
  selector:
    matchLabels:
      app: database
  replicas: 2
  serviceName: mysql-svc
  template:
    metadata:
      labels:
        app: database
    spec:
      terminationGracePeriodSeconds: 10
      containers:
        - name: dbserver
          image: registry.ocp4.example.com:8443/redhattraining/mysql-app:v1
          ports:
            - name: database
              containerPort: 3306
          env:
            - name: MYSQL_USER
              value: "redhat"
            - name: MYSQL_PASSWORD
              value: "redhat123"
            - name: MYSQL_DATABASE
              value: "sakila"
          volumeMounts:
            - name: data
              mountPath: /var/lib/mysql
  volumeClaimTemplates:
    - metadata:
        name: data
      spec:
        accessModes: [ "ReadWriteOnce" ]
        storageClassName: "lvms-vg1"
        resources:
          requests:
            storage: 1Gi
```

Edit the /home/student/DO180/labs/storage-statefulsets/service-db.yaml file and modify the placeholder text to match the following content:

```
apiVersion: v1
kind: Service
metadata:
  name: mysql-svc
  labels:
    app: database
spec:
  ports:
    - port: 3306
      name: mysql
  clusterIP: None
  selector:
    app: database
```

Create the mysql-svc headless service.

```
[student@workstation ~]$ oc create -f \
/home/student/DO180/labs/storage-statefulsets/service-db.yml
service/mysql-svc created
```

Create the dbserver stateful set.

```
[student@workstation ~]$ oc create -f \
/home/student/DO180/labs/storage-statefulsets/statefulset-db.yml
statefulset.apps/bdserver created
```

Run the watch command and wait until the two stateful set pods are running.

```
[student@workstation ~]$ watch oc get pod \
-l app=database -o wide
NAME      READY   STATUS    ...   AGE     IP
dbserver-0 1/1     Running   ...  10s    10.8.0.112
dbserver-1 1/1     Running   ...  5s     10.8.0.113
```

Press **Ctrl+C** to end the watch command after the pods display Running in the STATUS column.

List the mysql-svc service endpoints, which correspond to the pods' IP addresses.

```
[student@workstation ~]$ oc get endpoints
NAME      ENDPOINTS          AGE
mysql-svc  10.8.0.112:3306,10.8.0.113:3306  10m
```

The endpoint IPs are different on your system.

Connect to the dbserver-1 pod by using its unique DNS name.

```
[student@workstation ~]$ oc rsh dbserver-1 \
curl -v telnet://dbserver-1.mysql-svc:3306
* Rebuilt URL to: telnet://dbserver-1.mysql-svc:3306/
* Trying 10.8.0.113...
* TCP_NODELAY set
* Connected to dbserver-1.mysql-svc (10.8.0.113) port 3306 (#0)
...output omitted...
```

The DNS name resolves IP that corresponds to the dbserver-1 pod that is listed in the previous command.

7. Confirm that each instance from the dbserver stateful set has a persistent volume claim. Add data to each stateful set pod.

Verify that the two persistent volume claims display Bound in the STATUS column.

```
[student@workstation ~]$ oc get pvc -l app=database
NAME      STATUS    ...   CAPACITY   ACCESS MODE ...
data-dbserver-0  Bound    ...  1Gi        RWO       ...
data-dbserver-1  Bound    ...  1Gi        RWO       ...
```

Verify that the dbserver-0 pod mounts the data-dbserver-0 PVC.

```
[student@workstation ~]$ oc get pod dbserver-0 -o json | \
jq .spec.volumes[0].persistentVolumeClaim.claimName
"data-dbserver-0"
```

Verify that the dbserver-1 pod mounts the data-dbserver-1 PVC.

```
[student@workstation ~]$ oc get pod dbserver-1 -o json | \
jq .spec.volumes[0].persistentVolumeClaim.claimName
"data-dbserver-1"
```

Create a table named items within the sakila database in the dbserver-0 pod.

```
[student@workstation ~]$ oc exec -it pod/dbserver-0 -- /bin/bash -c \
"mysql -uredhat -predhat123 sakila -e 'create table items (count INT);'"
mysql: [Warning] Using a password on the command line interface can be insecure.
```

Create a table named inventory within the sakila database in the dbserver-1 pod.

```
[student@workstation ~]$ oc exec -it pod/dbserver-1 -- /bin/bash -c \
"mysql -uredhat -predhat123 sakila -e 'create table inventory (count INT);'"
mysql: [Warning] Using a password on the command line interface can be insecure.
```

8. No data replication is configured at the application level for the dbserver stateful set. Verify that each database pod contains unique data.

Confirm that the dbserver-0 pod has a database named sakila with one table named items.

```
[student@workstation ~]$ oc exec -it pod/dbserver-0 -- /bin/bash -c \
"mysql -uredhat -predhat123 sakila -e 'show tables;'"
mysql: [Warning] Using a password on the command line interface can be insecure.
+-----+
| Tables_in_sakila |
+-----+
| items           |
+-----+
```

Confirm that the dbserver-1 pod has a database named sakila with one table named inventory.

```
[student@workstation ~]$ oc exec -it pod/dbserver-1 -- /bin/bash -c \
"mysql -uredhat -predhat123 sakila -e 'show tables;'"
mysql: [Warning] Using a password on the command line interface can be insecure.
+-----+
| Tables_in_sakila |
+-----+
| inventory        |
+-----+
```

9. Delete a pod in the dbserver stateful set. A new stateful set pod is created with the same ordinal index, and the pod mounts its corresponding PVC. The data remains in the PVC and it is accessible from the pod.

Delete the dbserver-0 pod in the dbserver stateful set.

```
[student@workstation ~]$ oc delete pod dbserver-0
pod "dbserver-0" deleted
```

Verify that a new dbserver-0 pod is created.

```
[student@workstation ~]$ oc get pods -l app=database
NAME      READY   STATUS    RESTARTS   AGE
dbserver-0 1/1     Running   0          4s
dbserver-1 1/1     Running   0          15m
```

Verify that the dbserver-0 pod mounts the data-dbserver-0 PVC.

```
[student@workstation ~]$ oc get pod dbserver-0 -o json | \
jq .spec.volumes[0].persistentVolumeClaim.claimName
"data-dbserver-0"
```

Verify that the new dbserver-0 pod has the items table within the sakila database.

```
[student@workstation ~]$ oc exec -it pod/dbserver-0 -- /bin/bash -c \
"mysql -uredhat -predhat123 sakila -e 'show tables;'"
mysql: [Warning] Using a password on the command line interface can be insecure.
+-----+
| Tables_in_sakila |
+-----+
| items           |
+-----+
```

Finish

On the workstation machine, use the lab command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish storage-statefulsets
```