

Guided Exercise: Application Autoscaling

Configure an autoscaler for an application and then load test that application to observe scaling up.

Outcomes

- Deploy an application.
- Manually scale the application.
- Configure application resource requests.
- Deploy the application with a Horizontal Pod Autoscaler.

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the cluster is accessible and that all resources are available for this exercise. It also creates the `reliability-autoscaling` project.

```
[student@workstation ~]$ lab start reliability-autoscaling
```

Instructions

1. Log in to the OpenShift cluster as the developer user with developer as the password. Use the `reliability-autoscaling` project.

Log in to the OpenShift cluster.

```
[student@workstation ~]$ oc login -u developer -p developer \
  https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

Set the `reliability-autoscaling` project as the active project.

```
[student@workstation ~]$ oc project reliability-autoscaling
...output omitted...
```

2. Create the `loadtest` deployment, service, and route. The deployment uses the `registry.ocp4.example.com:8443/redhat-training/loadtest:v1.0` container image, which provides a web application. The web application exposes an API endpoint that creates a CPU-intensive task when queried.

Review the `~/DO180/labs/reliability-autoscaling/loadtest.yml` resource file that the `lab` command prepared. The container specification does not include the resources section that you use to specify CPU requests and limits. You configure that section in another step. Do not change the file for now.

```
apiVersion: v1
kind: List
metadata: {}
items:
  - apiVersion: apps/v1
    kind: Deployment
    ...output omitted...
    spec:
      containers:
        - image: registry.ocp4.example.com:8443/redhat-training/loadtest:v1.0
          name: loadtest
          readinessProbe:
            failureThreshold: 3
            httpGet:
              path: /api/loadtest/v1/healthz
              port: 8080
              scheme: HTTP
            periodSeconds: 10
            successThreshold: 1
            timeoutSeconds: 1

        - apiVersion: v1
          kind: Service
        ...output omitted...

  - apiVersion: route.openshift.io/v1
    kind: Route
    ...output omitted...
```

Use the `oc apply` command to create the application.

```
[student@workstation ~]$ oc apply -f \
~/DO180/labs/reliability-autoscaling/loadtest.yml
deployment.apps/loadtest created
service/loadtest created
route.route.openshift.io/loadtest created
```

Run the `watch` command and wait for the pod to report the Running status. The name of the pod is different in your cluster.

```
[student@workstation ~]$ watch oc get pods
Every 2.0s: oc get pods

NAME          READY   STATUS    RESTARTS   AGE
loadtest-86d987b9bf-w66x9  1/1     Running   0          40s
```

Press **Ctrl+C** to end the `watch` command.

- Configure a horizontal pod autoscaler resource for the `loadtest` deployment. Set the minimum number of replicas to 2 and the maximum to 20. Set the average CPU usage to 50% of the CPU requests attribute.

The horizontal pod autoscaler does not work, because the `loadtest` deployment does not specify requests for CPU usage.

Use the `oc autoscale` command to create the horizontal pod autoscaler resource.

```
[student@workstation ~]$ oc autoscale deployment/loadtest --min 2 --max 20 \
--cpu-percent 50
horizontalpodautoscaler.autoscaling/loadtest autoscaled
```

Retrieve the status of the `loadtest` horizontal pod autoscaler resource. The unknown value in the `TARGETS` column indicates that OpenShift cannot compute the current CPU usage of the `loadtest` deployment. The deployment must include the CPU requests attribute for OpenShift to be able to compute the CPU usage.

```
[student@workstation ~]$ oc get hpa loadtest
NAME      REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
loadtest  Deployment/loadtest  <unknown>/50%  2          20         2          74s
```

Get more details about the resource status. You might have to rerun the command several times. Wait up to three minutes for the command to report the warning message.

```
[student@workstation ~]$ oc describe hpa loadtest
Name:                                     loadtest
Namespace:                                reliability-autoscaling
...output omitted...
Conditions:
  Type      Status  Reason           Message
  ----      ----   ----            -----
  AbleToScale  True    SucceededGetScale  the HPA controller was able to get the target's current scale
  ScalingActive False   FailedGetResourceMetric the HPA was unable to compute the replica count: failed to get cpu utilization: missing request for cpu
Events:
  Type      ... Message
  ----      ...
...output omitted...
  Warning   ... failed to get cpu utilization: missing request for cpu
...output omitted...
```

Delete the horizontal pod autoscaler resource. You re-create the resource in another step, after you fix the `loadtest` deployment.

```
[student@workstation ~]$ oc delete hpa loadtest
horizontalpodautoscaler.autoscaling "loadtest" deleted
```

Delete the `loadtest` application.

```
[student@workstation ~]$ oc delete -f \
~/DO180/labs/reliability-autoscaling/loadtest.yml
deployment.apps "loadtest" deleted
service "loadtest" deleted
route.route.openshift.io "loadtest" deleted
```

- Add a CPU resource section to the `~/DO180/labs/reliability-autoscaling/loadtest.yml` file. Redeploy the application from the file.

Edit the `~/DO180/labs/reliability-autoscaling/loadtest.yml` file, and configure the CPU limits and requests for the loadtest deployment. The pod needs 25 millicores to operate, and must not consume more than 100 millicores.

You can compare your work with the completed `~/DO180/solutions/reliability-autoscaling/loadtest.yml` file that the lab command prepared.

```
...output omitted...
spec:
  containers:
    - image: registry.ocp4.example.com:8443/redhattraining/loadtest:v1.0
      name: loadtest
      readinessProbe:
        failureThreshold: 3
        httpGet:
          path: /api/loadtest/v1/healthz
          port: 8080
          scheme: HTTP
        periodSeconds: 10
        successThreshold: 1
        timeoutSeconds: 1
      resources:
        requests:
          cpu: 25m
        limits:
          cpu: 100m
...output omitted...
```

Use the `oc apply` command to deploy the application from the file.

```
[student@workstation ~]$ oc apply -f \
~/DO180/labs/reliability-autoscaling/loadtest.yml
deployment.apps/loadtest created
service/loadtest created
route.route.openshift.io/loadtest created
```

Run the `watch` command and wait until the pod is running. The name of the pod is different in your cluster.

```
[student@workstation ~]$ watch oc get pods
NAME           READY   STATUS    RESTARTS   AGE
loadtest-667bdcdc99-vhc9x  1/1     Running   0          36s
```

Press **Ctrl+C** to end the `watch` command after the pod displays `Running` in the `STATUS` column.

- Manually scale the `loadtest` deployment by first increasing and then decreasing the number of running pods.

Scale up the `loadtest` deployment to five pods.

```
[student@workstation ~]$ oc scale deployment/loadtest --replicas 5
deployment.apps/loadtest scaled
```

Run the `watch` command and wait until the five pods are running. The names of the pods are different in your cluster.

```
[student@workstation ~]$ watch oc get pods
NAME           READY   STATUS    RESTARTS   AGE
loadtest-667bdcdc99-5fcvh  1/1     Running   0          43s
loadtest-667bdcdc99-dpspr  1/1     Running   0          42s
loadtest-667bdcdc99-hkssk  1/1     Running   0          43s
loadtest-667bdcdc99-vhc9x  1/1     Running   0          8m11s
loadtest-667bdcdc99-z5n9q  1/1     Running   0          43s
```

Press **Ctrl+C** to end the `watch` command after the pods display `Running` in the `STATUS` column.

Scale down the `loadtest` deployment back to one pod.

```
[student@workstation ~]$ oc scale deployment/loadtest --replicas 1
deployment.apps/loadtest scaled
```

Run the `watch` command and wait until the pod is running. The name of the pod is different in your cluster.

```
[student@workstation ~]$ watch oc get pods
NAME           READY   STATUS    RESTARTS   AGE
loadtest-667bdcdc99-vhc9x  1/1     Running   0          11m
```

Press **Ctrl+C** to end the watch command after the pod displays Running in the STATUS column.

- Configure a horizontal pod autoscaler resource for the `loadtest` deployment. Set the minimum number of replicas to 2 and the maximum to 20. Set the average CPU usage to 50% of the CPU request attribute.

Use the `oc autoscale` command to create the horizontal pod autoscaler resource.

```
[student@workstation ~]$ oc autoscale deployment/loadtest --min 2 --max 20 \
--cpu-percent 50
horizontalpodautoscaler.autoscaling/loadtest autoscaled
```

Open a new terminal window and run the `watch` command to monitor the `oc get hpa loadtest` command. Wait up to five minutes for the `loadtest` horizontal pod autoscaler to report usage in the `TARGETS` column.

Notice that the horizontal pod autoscaler scales up the deployment to two replicas, to conform with the minimum number of pods that you configured.

```
[student@workstation ~]$ watch oc get hpa loadtest
Every 2.0s: oc get hpa loadtest                         workstation: Fri Mar  3 06:26:24 2023
NAME      REFERENCE          TARGETS    MINPODS   MAXPODS   REPLICAS   AGE
loadtest  Deployment/loadtest  0%/50%     2          20         2          52s
```

Leave the command running, and do not interrupt it.

- Increase the CPU usage by sending requests to the `loadtest` application API.

Use the `oc get route` command to retrieve the URL of the application.

```
[student@workstation ~]$ oc get route loadtest
NAME      HOST/PORT
loadtest  loadtest-reliability-autoscaling.apps.ocp4.example.com ...
```

Send a request to the application API to simulate additional CPU pressure on the container. Do not wait for the `curl` command to complete, and continue with the exercise. After one minute, the command reports a timeout error that you can ignore.

```
[student@workstation ~]$ curl \
loadtest-reliability-autoscaling.apps.ocp4.example.com/api/loadtest/v1/cpu/1
<html><body><h1>504 Gateway Time-out</h1>
The server didn't respond in time.
</body></html>
```

Watch the output of the `oc get hpa loadtest` command in the second terminal. After a minute, the horizontal pod autoscaler detects an increase in the CPU usage and deploys additional pods.

NOTE

The increased activity of the application does not immediately trigger the autoscaler. Wait a few moments if you do not see any changes to the number of replicas.

You might need to run the `curl` command multiple times before the application uses enough CPU to trigger the autoscaler.

The CPU usage and the number of replicas on your system probably differ.

```
Every 2.0s: oc get hpa loadtest                         workstation: Fri Mar  3 07:20:19 2023
NAME      REFERENCE          TARGETS    MINPODS   MAXPODS   REPLICAS   AGE
loadtest  Deployment/loadtest  220%/50%   2          20         9          16m
```

Wait five minutes after the `curl` command completes. The `oc get hpa loadtest` command shows that the CPU load decreases.

NOTE

Although the horizontal pod autoscaler resource can be quick to scale up, it is slower to scale down.

```
Every 2.0s: oc get hpa loadtest          workstation: Fri Mar  3 07:23:11 2023
NAME      REFERENCE      TARGETS      MINPODS      MAXPODS      REPLICAS      AGE
loadtest  Deployment/loadtest  0%/50%        2            20            9           18m
```

Optional: Wait for the `loadtest` application to scale down. It takes five additional minutes for the horizontal pod autoscaler to scale down to two replicas.

```
Every 2.0s: oc get hpa loadtest          workstation: Fri Mar  3 07:29:12 2023
NAME      REFERENCE      TARGETS      MINPODS      MAXPODS      REPLICAS      AGE
loadtest  Deployment/loadtest  0%/50%        2            20            2           24m
```

Press **Ctrl+C** to quit the `watch` command. Close that second terminal when done.

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish reliability-autoscaling
```