

Quiz: Configure an Application for Reliability

Configure an Application for Reliability

Choose the correct answers to the following questions:

1.

In the previous lab, given your node capacity and the application's initial memory request, how many pods can OpenShift schedule?

- A The request is set to 8 GiB; therefore, no pod can be scheduled.
- B The request is set to 8 GiB; therefore, only one pod can be scheduled.
- C The request is set to 512 MiB; therefore, no pod can be scheduled.
- D The request is set to 512 MiB; therefore, three pods can be scheduled.

2.

In the previous lab, which deployment setting is incorrect and prevents the application from starting?

- A The readiness probe is misconfigured and always fails.
- B The liveness probe is misconfigured and always fails.
- C The memory request is too high. No compute nodes have enough memory to accommodate the application.
- D The memory limit is too low. The OOM kernel subsystem keeps killing the container processes.

3.

In the previous lab, which setting is initially absent from the deployment, but does not prevent the application from running?

- A The liveness probe is not defined.
- B The readiness probe is not defined.

- C The memory limit is not set.
- D The CPU limit is not set.

4.

In the previous lab, which metric and target value do you configure in the HorizontalPodAutoscaler resource to trigger automatic scaling for the `longload` deployment?

- A A minimum of one and a maximum of three replicas.
- B A fixed memory usage of 480 MiB.
- C An average memory utilization of 60%.
- D An average CPU utilization of 60%.